

Zafeer Rangoonwala

US Work Authorized: F-1 STEM OPT (No sponsorship required for 3 years)

Dallas, TX | +1 (469) 268-0532 | zafeer02@gmail.com | linkedin.com/in/zafeerr/ | github.com/Zafeer-R | zafeerrangoonwala.com

SUMMARY

Software Engineer and Data/AI Engineer with experience building real-time data pipelines, AI-powered applications, and cloud analytics systems. Skilled in Python, SQL, Node.js, PySpark, Kafka, Docker, and LLM/RAG integration. **Hackathon Winner** and MS CS graduate from **University of Texas at Dallas**.

EDUCATION

The University of Texas at Dallas, Texas (CGPA 3.5/4.0)	USA
Master's, Computer Science	Aug 2024 - May 2026
Indus University, India (CGPA 3.59/4.0)	INDIA
Bachelor's, Computer Science	Oct 2020 - May 2024

EXPERIENCE

The University of Texas at Dallas, USA: Technology and Operations Worker May 2025 - May 2026

- Sole IT support for 40+ member international centers, managing all hardware, access, and device provisioning end-to-end.
- Owned full asset lifecycle — imaging, onboarding, deprovisioning, and issue resolution — streamlining workflows and eliminating support bottlenecks.

Delta Systems, Ahmedabad, India: Data Scientist Jan 2024 - Jul 2024

- Engineered ML workflows using Random Forest, XGBoost, and time-series forecasting on sensor data, improving prediction accuracy by 23%.
- Built automated Python and SQL ingestion and validation pipelines to clean and monitor 1M+ daily production and equipment records.
- Designed anomaly detection solutions to identify valve failures and pressure deviations, visualizing insights via Tableau and Power BI.
- Deployed ML-driven anomaly detection in Agile teams, cutting troubleshooting time by 70% and deployment cycles from 4 to 2 weeks
- Implemented k-means clustering with scikit-learn on customer data, creating segments that supported targeted marketing strategies which resulted in 10% better throughput.

KnockOnce, Remote: Software Engineer Intern Jan 2023 - Mar 2023

- Developed a high-traffic platform (sumosave.in) from the ground up using provided designs, building responsive pages serving nearly **10,000 users**.
- Built a scalable backend architecture using MySQL, designing normalized database schemas for partnerships, careers, and admin management while implementing admin-side CRUD operations for dynamic job posting and content management.

SKILLS

Languages	Python, SQL, Java, JavaScript, C/C++, PySpark
Frontend:	React, TypeScript, JavaScript (ES6+), HTML5, CSS3, Tailwind CSS, Vite, Webpack
Backend:	Node.js, Express, FastAPI, Django, Flask, Spring Boot, REST APIs, GraphQL, WebSockets, SSE Streaming
Databases:	Supabase (PostgreSQL), Firebase, MongoDB, MySQL
AI/GenAI:	LLMs, OpenAI GPT, Claude, Gemini, Groq/Llama 3, LangChain, RAG, MCP, Prompt Engineering, Fine-Tuning, n8n Workflow Automation
Cloud & DevOps:	AWS (SageMaker, EC2, S3, Lambda, CloudWatch), GCP (Vertex AI, BigQuery), Docker, Kubernetes, Terraform, CI/CD, GitHub Actions
Tools	Git, Jupyter, Databricks, Tableau, Power BI, Postman, Apache Kafka, Kibana, ElevenLabs
Certifications:	Anthropic Academy: Claude 101, Building with Claude API, MCP, Claude Code, Agent Skills

PROJECTS (zafeerrangoonwala.com/#projects)

RealPage AI | Agentic RAG Property Intelligence System | [GitHub](#)

Built a production-grade RAG system using LangGraph where an LLM planner intelligently selects between SQL, graph, and vector retrievals, and Groq synthesizes factual answers with persistent conversation memory, demonstrating agentic routing, multi-modal retrieval, hallucination mitigation through prompt engineering, and streaming API design. **Tech:** Python, FastAPI, LangGraph, Groq, PostgreSQL, Neo4j, Qdrant, React, TypeScript, LangSmith

GeoSearch | Search Engine For Geology | [GitHub](#)

Built GeoSearch, a geology-focused search engine that crawls earth-science web content, builds searchable indexes, supports multiple ranking models, and enhances retrieval with query expansion and clustering. Developed an end-to-end IR pipeline with a React/FastAPI interface for comparing relevance models, clustered results, expanded queries, and external search engines. **Tech:** Python, FastAPI, React, Vite, JavaScript, scikit-learn, NumPy, SciPy, NLTK, PyStemmer, Uvicorn, Pydantic, Requests, SerpAPI, TF-IDF, BM25, PageRank, HITS, MiniBatchKMeans, Agglomerative Clustering

SignalTrace | Real-Time Financial News Analyser | [GitHub](#)

Built a real-time financial news intelligence pipeline using Kafka, PySpark Structured Streaming, spaCy, Elasticsearch, and Kibana to ingest live news, extract named entities, compute rolling mention trends, and visualize top entities through live dashboards, demonstrating event-driven architecture, stream processing, and real-time analytics. **Tech:** Python, Apache Kafka, PySpark Structured Streaming, spaCy, Elasticsearch, Kibana, Logstash, FintHub API

AutoTracker | Scalable Vehicle Listing Event Pipeline | [GitHub](#)

Built an event-driven vehicle listing platform using Spring Boot, Kafka, MySQL, and React, enabling asynchronous processing of listing events (create, price update, sold), DLQ-based error handling, REST APIs, and real-time listing history visualization to demonstrate distributed architecture and event streaming. **Tech:** Java, Spring Boot, Spring Kafka, Spring Data JPA, Hibernate, MySQL, Apache Kafka, Docker, Maven, JUnit, Mockito, React, TypeScript, Vite, Docker